

## Practical 1: Database Design and Table Creation

### Aim

To design a relational database schema for a college management system and implement it by creating the required tables using SQL CREATE TABLE command.

### Theory

A relational database is a collection of data items with pre-defined relationships between them. These items are organized as a set of tables with columns and rows. The design process involves identifying entities (e.g., Student, Department, Course), their attributes, and the relationships between them.

The SQL CREATE TABLE command is a Data Definition Language (DDL) command used to define and create a new table in the database.

### Syntax:

```
CREATE TABLE table_name ( column1 datatype  
[constraints], column2 datatype [constraints],  
column3 datatype [constraints],  
...  
[table_constraints]  
);
```

### Program:

```
-- Create the college database
```

```
CREATE DATABASE CollegeDB;
```

```
-- Select the database to use
```

```
USE CollegeDB;
```

```
-- Create the Departments table first (since Students and Courses depend on it)
```

```
CREATE TABLE Departments (  
DeptID INT PRIMARY KEY AUTO_INCREMENT,  
DeptName VARCHAR(50) NOT NULL,  
HOD VARCHAR(50)  
);
```

```
-- Create the Students table
```

```
CREATE TABLE Students (  
StudentID INT PRIMARY KEY AUTO_INCREMENT,  
FirstName VARCHAR(50) NOT NULL,  
LastName VARCHAR(50),  
DateOfBirth DATE,  
Email VARCHAR(100) UNIQUE,  
DeptID INT,  
FOREIGN KEY (DeptID) REFERENCES Departments(DeptID)  
    ON DELETE SET NULL  
    ON UPDATE CASCADE  
);
```

```
-- Create the Courses table
```

```
CREATE TABLE Courses (  
CourseID INT PRIMARY KEY AUTO_INCREMENT,  
CourseName VARCHAR(100) NOT NULL,
```

```
Credits INT CHECK (Credits > 0),  
DeptID INT,  
FOREIGN KEY (DeptID) REFERENCES Departments(DeptID)  
    ON DELETE CASCADE  
    ON UPDATE CASCADE  
);
```

**Sample Output (Execution result)**

When you run these commands in MySQL Workbench or any SQL terminal, you'll see:

```
Query OK, 1 row affected (0.01 sec)  
Database changed  
Query OK, 0 rows affected (0.02 sec)  
Query OK, 0 rows affected (0.01 sec)  
Query OK, 0 rows affected (0.01 sec)
```

---

**Conclusion :**We successfully designed a relational database schema for a college and implemented it by creating three related tables (Departments, Students, Courses) using the CREATE TABLE command.

## Practical 2: Applying Constraints

### Aim

To apply various integrity constraints to the tables, such as PRIMARY KEY, FOREIGN KEY, NOT NULL, UNIQUE, CHECK, and DEFAULT.

### Theory

Integrity constraints are rules that are applied to columns to ensure the accuracy and reliability of the data. They are a fundamental part of good database design.

- **PRIMARY KEY:** Uniquely identifies each row in a table.
- **FOREIGN KEY:** Establishes a link between data in two tables.
- **NOT NULL:** Ensures that a column cannot have a NULL value.
- **UNIQUE:** Ensures that all values in a column are different.
- **CHECK:** Ensures that a value in a column meets a specific condition.
- **DEFAULT:** Provides a default value for a column when none is specified.

### Program

-- Assuming tables from Practical 1 are already created

```
USE CollegeDB;
```

-- Add a UNIQUE constraint to the Email column in Students table (already done in Pr1)

-- This line is safe; if already unique, it will throw a warning if re-added

```
ALTER TABLE Students
```

```
ADD CONSTRAINT UQ_Email UNIQUE (Email);
```

-- Add a CHECK constraint to the Students table to ensure age is above 18

-- MySQL supports CHECK constraints from version 8.0.16+

```
ALTER TABLE Students
```

```
ADD CONSTRAINT CHK_Age CHECK (TIMESTAMPDIFF(YEAR, DateOfBirth, CURDATE()) >= 18);
```

-- Add a DEFAULT value for the HOD column in Departments table

```
ALTER TABLE Departments
```

```
ALTER COLUMN HOD SET DEFAULT 'To be assigned';
```

-- Add a new table 'Faculty' with constraints

```
CREATE TABLE Faculty (
    FacultyID INT PRIMARY KEY AUTO_INCREMENT,
    FullName VARCHAR(100) NOT NULL,
    DeptID INT,
    Salary DECIMAL(10, 2) CHECK (Salary >= 30000),
    CONSTRAINT FK_FacultyDept FOREIGN KEY (DeptID)
        REFERENCES Departments(DeptID)
        ON DELETE SET NULL
        ON UPDATE CASCADE
);
```

### Output:

Database changed

Query OK, 0 rows affected (0.02 sec)

Query OK, 0 rows affected (0.01 sec)

Query OK, 0 rows affected (0.01 sec)

Query OK, 0 rows affected (0.02 sec)

Query OK, 0 rows affected (0.03 sec)

## **Conclusion**

We successfully applied various integrity constraints to the tables in our college database. This ensures data consistency and prevents invalid or incomplete data from being entered.

## Practical 3: ALTER, UPDATE, and DELETE SQL statements

### Aim

To use the ALTER, UPDATE, and DELETE commands to modify the structure and data of existing tables.

### Theory

These are essential Data Manipulation Language (DML) and Data Definition Language (DDL) commands for maintaining a database.

- **ALTER TABLE:** Used to add, delete, or modify columns in an existing table. It can also be used to add or drop constraints.
- **UPDATE:** Used to modify existing records in a table. The WHERE clause is crucial to specify which records to update.
- **DELETE:** Used to delete existing records from a table. The WHERE clause is crucial to specify which records to delete.

### Program

-- Assuming tables from previous practicals already exist

```
USE CollegeDB;
```

-- Add a new column 'Semester' to the Students table

```
ALTER TABLE Students
```

```
ADD COLUMN Semester INT DEFAULT 1;
```

-- Update the Semester for all existing students (not necessary if default is set, but included)

```
UPDATE Students
```

```
SET Semester = 1;
```

-- Add some sample data to work with

-- (Avoid manually setting IDs if AUTO\_INCREMENT is enabled)

```
INSERT INTO Departments (DeptName, HOD) VALUES ('Computer Science', 'Dr. Rao');
```

```
INSERT INTO Departments (DeptName, HOD) VALUES ('Mechanical Engineering', 'Dr. Verma');
```

-- Check the DeptIDs assigned automatically

```
SELECT * FROM Departments;
```

-- Insert sample students (using auto-generated StudentID)

```
INSERT INTO Students (FirstName, LastName, DateOfBirth, Email, DeptID)
```

```
VALUES
```

```
('John', 'Doe', '2003-05-12', 'john.doe@gmail.com', 1),
```

```
('Jane', 'Smith', '2002-08-20', 'jane.smith@gmail.com', 2);
```

-- Update a specific student's department

```
UPDATE Students
```

```
SET DeptID = 1
```

```
WHERE FirstName = 'Jane' AND LastName = 'Smith';
```

-- Delete a specific department

--  This will fail if foreign key constraint prevents deletion.

-- To allow it, ensure ON DELETE CASCADE or SET NULL is set in Students table.

```
DELETE FROM Departments
```

```
WHERE DeptID = 2;
```

-- Verify final data

```
SELECT * FROM Departments;
```

```
SELECT * FROM Students;
```

**Output:**

**Expected Outputs**

After all operations, your tables will look like:

**Departments**

DeptID	DeptName	HOD
1	Computer Science	Dr. Rao

**Students**

StudentID	FirstName	LastName	DateOfBirth	Email	DeptID	Semester
1	John	Doe	2003-05-12	<a href="mailto:john.doe@gmail.com">john.doe@gmail.com</a>	1	1
2	Jane	Smith	2002-08-20	<a href="mailto:jane.smith@gmail.com">jane.smith@gmail.com</a>	1	1

---

**Conclusion :**

We successfully used ALTER TABLE to modify a table's structure, UPDATE to change data within the tables, and DELETE to remove records. These commands are essential for database maintenance.

## Practical 4: Joins

### Aim

To write queries to retrieve data from multiple tables using different types of joins.

### Theory

A JOIN is a clause that combines rows from two or more tables based on a related column between them.

- **INNER JOIN:** Returns rows when there is at least one match in both tables.
- **LEFT JOIN:** Returns all rows from the left table, and the matched rows from the right table. The result is NULL from the right side if there is no match.
- **RIGHT JOIN:** Returns all rows from the right table, and the matched rows from the left table. The result is NULL from the left side when there is no match.
- **FULL OUTER JOIN:** Returns all rows when there is a match in one of the tables.

### Program

-- Inner Join: Get students with their department names

```
SELECT S.FirstName, S.LastName, D.DeptName  
FROM Students AS S  
INNER JOIN Departments AS D ON S.DeptID = D.DeptID;
```

-- Left Join: Get all students and their department names (if they have one)

```
SELECT S.FirstName, S.LastName, D.DeptName  
FROM Students AS S  
LEFT JOIN Departments AS D ON S.DeptID = D.DeptID;
```

-- Right Join: Get all departments and their student names (if they have any)

```
SELECT S.FirstName, S.LastName, D.DeptName  
FROM Students AS S  
RIGHT JOIN Departments AS D ON S.DeptID = D.DeptID;
```

### Output:

#### Inner join:

FirstName	LastName	DeptName
John	Doe	Computer Science
Jane	Smith	Computer Science

#### Left join:

FirstName	LastName	DeptName
John	Doe	Computer Science
Jane	Smith	Computer Science
Ramesh	Patil	NULL

#### Right join:

FirstName	LastName	DeptName
John	Doe	Computer Science
Jane	Smith	Computer Science
NULL	NULL	Mechanical Engineering

## **Conclusion**

We successfully used INNER, LEFT, and RIGHT joins to retrieve combined data from the Students and Departments tables. Joins are fundamental for working with relational databases.

## Practical 5: Aggregate Functions

### Aim

To write queries that implement aggregate functions like MAX(), MIN(), AVG(), and COUNT().

### Theory

Aggregate functions perform a calculation on a set of values and return a single value. They are often used with the GROUP BY clause.

- **COUNT()**: Returns the number of rows that match a specified criterion.
- **SUM()**: Returns the total sum of a numeric column.
- **AVG()**: Returns the average value of a numeric column.
- **MIN()**: Returns the smallest value of the selected column.
- **MAX()**: Returns the largest value of the selected column.

### Program

-

-- Count the number of students

```
SELECT COUNT(*) FROM Students;
```

-- Find the maximum and minimum number of credits for courses

```
SELECT MAX(Credits), MIN(Credits) FROM Courses;
```

-- Calculate the average salary of faculty

```
SELECT AVG(Salary) FROM Faculty;
```

-- Count the number of students in each department

```
SELECT DeptID, COUNT(StudentID)  
FROM Students  
GROUP BY DeptID;
```

### Conclusion :

We successfully used a variety of aggregate functions to perform calculations on our data, such as counting students, finding max/min values, and calculating averages. These functions are essential for data analysis.

## Practical 6: Subqueries and Nested Queries

### Aim

To implement the concept of nested queries and subqueries to retrieve data that depends on the result of another query.

### Theory

A subquery, or nested query, is a query that is embedded within another query. The inner query executes first, and its result is used by the outer query.

- **Subquery in WHERE clause:** Used to filter data based on the result of the inner query. ● **Subquery in FROM clause:** The result of the subquery is treated as a temporary table.
- **Subquery in SELECT clause:** Used to return a single value for each row of the outer query.

### Program

-- Find students in a specific department using a subquery

```
SELECT FirstName, LastName  
FROM Students  
WHERE DeptID = (SELECT DeptID FROM Departments WHERE DeptName = 'Computer  
Science');
```

-- Find the courses that have more credits than the average credits

```
SELECT CourseName, Credits  
FROM Courses  
WHERE Credits > (SELECT AVG(Credits) FROM Courses);
```

-- Find the departments that have at least 2 students

```
SELECT DeptName  
FROM Departments  
WHERE DeptID IN (SELECT DeptID FROM Students GROUP BY DeptID HAVING COUNT(*) >=  
2);
```

### Conclusion

We successfully used subqueries to perform more complex data retrieval. This demonstrates how to combine the power of multiple queries to achieve a desired result.

## Practical 7: Views, Sequences, and Indices

### Aim

To create and manage views, sequences, and indices to enhance database performance and security.

### Theory

- **Views:** A virtual table based on the result-set of an SQL statement. It does not store data but rather represents a predefined query. Views are used for security and to simplify complex queries.
- **Sequences:** A database object that generates a unique number for each request. They are often used for automatically generating primary key values.
- **Indices:** A data structure that improves the speed of data retrieval operations on a database table. An index is a pointer to the data in the table.

### Program

```
-- Create a view to show student names and department names  
CREATE VIEW StudentDeptView AS  
SELECT S.FirstName, S.LastName, D.DeptName  
FROM Students AS S  
INNER JOIN Departments AS D ON S.DeptID = D.DeptID;
```

```
-- Query the new view  
SELECT * FROM StudentDeptView;
```

```
-- Create an index on the LastName column of the Students table  
CREATE INDEX idx_student_lastname ON Students(LastName);
```

```
-- Create a sequence for student IDs  
CREATE SEQUENCE StudentID_Seq  
START WITH 100  
INCREMENT BY 1;
```

```
-- Use the sequence to insert a new student  
INSERT INTO Students (StudentID, FirstName, LastName, DeptID)  
VALUES (NEXT VALUE FOR StudentID_Seq, 'Alice', 'Smith', 101);
```

### Conclusion :

We successfully created a view for simplified data access, an index to optimize query performance, and a sequence for automatic ID generation. These objects are crucial for building efficient and maintainable database applications.

## Practical 8: Triggers

### Aim

To perform queries for triggers.

### Theory

A trigger is a set of SQL statements that is executed automatically in response to certain events on a table, such as an INSERT, UPDATE, or DELETE operation. Triggers are used to enforce business rules, audit data changes, or automate tasks.

- **Syntax:**

```
CREATE TRIGGER trigger_name  
[BEFORE | AFTER] [INSERT | UPDATE | DELETE] ON table_name  
FOR EACH ROW  
BEGIN  
    -- SQL statements to be executed  
END;
```

- **BEFORE/AFTER:** Specifies when the trigger fires (before or after the event).
- **FOR EACH ROW:** The trigger will execute for every row affected by the DML statement.

### Program

```
-- Assuming a table 'StudentLog' exists to log changes
```

```
USE CollegeDB;
```

```
CREATE TABLE StudentLog (  
    LogID INT PRIMARY KEY AUTO_INCREMENT,    StudentID INT,  
    Operation VARCHAR(10),  
    OperationTime DATETIME  
);
```

```
-- Create a trigger that logs every student insertion
```

```
CREATE TRIGGER AfterInsertStudent  
AFTER INSERT ON Students  
FOR EACH ROW  
INSERT INTO StudentLog (StudentID, Operation, OperationTime)  
VALUES (NEW.StudentID, 'INSERT', NOW());
```

```
-- Insert a new student to test the trigger
```

```
INSERT INTO Students (StudentID, FirstName, LastName, DeptID) VALUES (10, 'Mark', 'Wilson',  
101);
```

```
-- Verify the log entry
```

```
SELECT * FROM StudentLog;
```

### Conclusion:

We successfully created a trigger that automatically logs every student insertion into a separate log table. This demonstrates how triggers can be used for automation and auditing.

## Practical 9: Insertion, Updation, and Deletion with Referential Integrity

### Aim

To perform operations for demonstrating the insertion, updation, and deletion using the referential integrity constraints.

### Theory

Referential integrity is a relational database concept that ensures that relationships between tables are consistent. It is enforced using the FOREIGN KEY constraint. When a foreign key exists, it must reference a valid row in the parent table. The ON DELETE and ON UPDATE clauses of a foreign key constraint define the behavior when a row in the parent table is modified.

- **ON DELETE CASCADE:** When a row in the parent table is deleted, the corresponding rows in the child table are also automatically deleted.
- **ON DELETE SET NULL:** When a row in the parent table is deleted, the foreign key column in the child table is set to NULL.
- **ON DELETE RESTRICT:** Prevents the deletion of a row in the parent table if it has corresponding rows in the child table. This is the default behavior.

### Program

```
-- First, drop the existing foreign key constraint
ALTER TABLE Students
DROP FOREIGN KEY Students_ibfk_1;

-- Now, recreate the foreign key with CASCADE behavior
ALTER TABLE Students
ADD CONSTRAINT FK_StudentDept CASCADE ON DELETE;

-- Insert some sample data to work with
INSERT INTO Departments (DeptID, DeptName) VALUES (103, 'Civil Engineering');
INSERT INTO Students (StudentID, FirstName, LastName, DeptID) VALUES (20, 'David', 'Lee', 103);
INSERT INTO Courses (CourseID, CourseName, DeptID) VALUES (301, 'Structural Analysis', 103);

-- Delete a department and observe the cascade effect
DELETE FROM Departments WHERE DeptID = 103;

-- Attempt to insert a student with a non-existent department ID (this should fail)
INSERT INTO Students (StudentID, FirstName, LastName, DeptID) VALUES (21, 'Emily', 'Clark', 999);
```

### Conclusion :

We demonstrated how referential integrity, enforced by foreign keys with CASCADE behavior, automatically handles related data when a record in the parent table is deleted. We also showed how the database prevents data inconsistency by rejecting an insertion that violates referential integrity.

## Practical 10: User and Role Management

### Aim

To create and manage users and their roles to control access to the database.

### Theory

User and role management is a critical aspect of database security.

- **CREATE USER:** Command to create a new user account.
- **GRANT:** Command to give privileges (permissions) to a user or role.
- **REVOKE:** Command to remove privileges from a user or role.
- **CREATE ROLE:** Command to create a role, which is a collection of privileges.

### Program

-- Assuming a root user

-- Create a new user

```
CREATE USER 'student_user'@'localhost' IDENTIFIED BY 'password123';
```

-- Create a role for faculty

```
CREATE ROLE 'faculty_role';
```

-- Grant privileges to the role

```
GRANT SELECT, INSERT, UPDATE ON CollegeDB.Faculty TO 'faculty_role';
```

-- Grant the role to the user

```
GRANT 'faculty_role' TO 'student_user'@'localhost';
```

-- Revoke the insert privilege from the user

```
REVOKE INSERT ON CollegeDB.Faculty FROM 'student_user'@'localhost';
```

-- Drop the user and role

```
DROP USER 'student_user'@'localhost';
```

```
DROP ROLE 'faculty_role';
```

### Conclusion

We successfully created a new user, a role, granted privileges to the role, and then assigned the role to the user. We also demonstrated how to revoke privileges. This shows how to implement basic access control for a database.